


Three orthogonal raw evidences

Null-CDF calibration percentile vs random pairs

Reliability weighting down-weight contradiction

Calibrated affinity



Each term is reported on the edge: the score is interpretable and auditable, never an opaque number.

Pure numpy, so the browser recomputes S live when the analyst reweights the evidences.